

# Deep Learning

**Instructor:** Tongyao Pang

**Notes Taker:** Zejin Lin

TSINGHUA UNIVERSITY.

linzj23@mails.tsinghua.edu.cn

[lzmjmaths.github.io](https://lzmjmaths.github.io)

March 7, 2025

## Contents

<b>1</b>	<b>Regression</b>	<b>2</b>
1.1	Binary classification problem . . . . .	2
1.2	Gradient Descent . . . . .	3
1.3	Adaptive learning rate . . . . .	3
<b>2</b>	<b>Vanishing or Exploding Gradient</b>	<b>4</b>
	<b>Index</b>	<b>6</b>
	<b>List of Theorems</b>	<b>7</b>

# 1 Regression

$$\min_{\omega \in \mathbb{R}^m} \frac{1}{2N} \|\Phi\omega - y\|^2 + \lambda C(\omega) \quad (1.1)$$

**Lasso:**  $C = \|\omega\|_1$ . **Ridge regression:**  $C = \|\omega\|_2$ .

**subgradient** of  $f$ :

$$\partial f(x_0) = \{g | f(x) \geq f(x_0) + g^T(x - x_0)\}$$

In particular,

$$\partial|x| = \begin{cases} 1, & x > 0 \\ -1, & x < 0 \\ [-1, 1], & x = 0 \end{cases}$$

## 1.1 Binary classification problem

**one-hot encoding** for the output  $\left\{\begin{pmatrix} 1 \\ 0 \end{pmatrix}, \begin{pmatrix} 0 \\ 1 \end{pmatrix}\right\}$ . It can be understood as the probability for each class and can take continuous values.

A **linear hypothesis space** is  $\{u(x) : u = \omega^T x, x \in \mathbb{R}^n, \omega \in \mathbb{R}^n\}$ .

**Softmax:** Map the extracted feature  $u$  to the space of one-hot codes

$$\mu = \frac{1}{1 + e^{-u}}, \quad 1 - \mu = \frac{e^{-u}}{1 + e^{-u}} = \frac{1}{1 + e^u}$$

$$KL(p, q) = \int p(\log p - \log q) \quad (1.2)$$

For  $p$  real probability, to minimize (1.2), suffices to minimize

$$-\int p \log q_\theta dx = -\sum_{x_i} \log q_\theta(x_i)$$

which is called **Maximum likelihood (cross entropy)**

$$-\sum \log p(y_i | x_i, \omega) = \sum -y_i \log \mu_i - (1 - y_i) \log(1 - \mu_i)$$

We reduce to minimize the thing above.

## 1.2 Gradient Descent

$$J(\theta) = \sum_{i=1}^N L(f_{\theta}(x_i), y_i), \quad \theta^{t+1} = \theta^t - \eta_t \frac{\partial J(\theta)}{\partial \theta} \Big|_{\theta=\theta^t}$$

For **empirical loss**

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N J_i(\theta), \quad J_i(\theta) = L(f_{\theta}(y_i), x_i)$$

**Stochastic Gradient Descent:** only compute gradients over a mini batch for each epoch

$$\theta_{k+1} = \theta_k - \eta \frac{1}{B} \sum_{i \in I_B} \nabla J_i(\theta_k)$$

$I_B$ , called **mini batch** are randomly sampled from the training data indexes.

The motivation to sample is that for distribution  $x, \frac{x_1 + \dots + x_N}{N}$  has the same mean  $\mu$  but less variance  $\frac{1}{N}\sigma^2$ , so in order to have more randomness and greatly reduce the computational cost, we choose less amount of data.

Randomness can help avoid getting stuck in local minimum.

SGD:  $x_{t+1} = x_t - \alpha \nabla f(x_t)$ .

SGD+**momentum**:

$$v_{t+1} = \rho v_t + \nabla f(X_t)$$

$$x_{t+1} = x_t - \alpha v_{t+1}$$

$v_t$  is the momentum which helps accelerate convergence by accumulating the gradients of past steps and smoothing out the oscillations, where  $\rho = 0.9$  or  $0.99$ .

## 1.3 Adaptive learning rate

$$r_t = r_{t-1} + \nabla f(x_t) \odot \nabla f(x_t)$$

$$x_{t+1} = x_t - \frac{\alpha}{\sqrt{r_t + \varepsilon}} \odot \nabla f(x_t)$$

For frequent features, the updates will be smaller, and for rare features, the updates will be larger.

**Notation**  $\odot$  is the multiplication for each component, which means:

$$\begin{pmatrix} a_1 \\ \vdots \\ a_n \end{pmatrix} \odot \begin{pmatrix} b_1 \\ \vdots \\ b_n \end{pmatrix} = \begin{pmatrix} a_1 b_1 \\ \vdots \\ a_n b_n \end{pmatrix}$$

**Exponential Moving Averaging(EMA)**

$$r_t = \beta r_{t-1} + (1 - \beta) \nabla f(x_t) \odot \nabla f(x_t)$$

$$x_{t+1} = x_t - \frac{\alpha}{\sqrt{r_t + \varepsilon}} \bigodot \nabla f(x_t)$$

RMSProp uses a moving average, avoiding overly aggressive decay compared with AdaGrad.

**Adaptive Moment Estimation(Adam):**RMASProp+Momentum

$$g_t = \nabla f(x_t)$$

$$v_t = \beta_1^t v_{t-1} + (1 - \beta_1^t) g_t, r_t = \beta_2^t r_{t-1} + (1 - \beta_2^t) g_t \odot g_t$$

$$v_t = \frac{v_t}{1 - \beta_1^t}, E_t = \frac{E_t}{1 - \beta_2^t}$$

$$x_{t+1} = x_t - \frac{\alpha}{\sqrt{r_t + \varepsilon}} \bigodot v_t$$

## 2 Vanishing or Exploding Gradient

The gradient of the Sigmoid function is very small most of the time, leading to vanishing gradients

We want to avoid exploding or vanishing in gradient.

**Sigmoid:**  $\sigma(z) = \frac{1}{1 + e^{-z}}.$

**Tanh:**  $\sigma(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}.$

**ReLU:**  $\text{ReLU}(z) = \begin{cases} z, & z > 0 \\ 0, & \text{otherwise} \end{cases}.$

**LeakyReLU:**  $\text{LeakyReLU}(z) = \begin{cases} z, & z > 0 \\ az, & \text{otherwise} \end{cases}.$  The gradient neither vanishes nor explodes;

it is computationally fast, but some neurons may not be activated.

LeakyReLU solves the issue with ReLU and is the most commonly used.

Consider the back propagation  $\frac{\partial J}{\partial x} = \frac{\partial J}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial J}{\partial x} W$ . Cumulate after multi-layers

$$\text{Var}\left(\frac{\partial J}{\partial x}\right) = \prod_i n_l \text{Var}(W_l) \text{Var}\left(\frac{\partial J}{\partial x_l}\right)$$

We want  $n_l \text{Var}(W_l) \sim 1$ .

After normalizing their variances, the updating becomes more steady and efficient.

To avoid variance becoming too small or too large in deep layers, normalize features in the network

$$\hat{x}_i = \frac{x_i - \mathbb{E}x_i}{\sqrt{\text{var}(x_i)}}$$

**Batch Normalization:** normalize features across samples within each batch.

# Index

Adaptive Moment Estimation(Adam), [4](#)

Batch Normalization, [5](#)

cross entropy, [2](#)

Exponential Moving Averaging(EMA), [4](#)

Lasso, [2](#)

LeakyReLU, [4](#)

linear hypothesis space, [2](#)

Maximum likelihood, [2](#)

mini batch, [3](#)

momentum, [3](#)

one-hot encoding, [2](#)

ReLU, [4](#)

Ridge regression, [2](#)

Sigmoid, [4](#)

Softmax, [2](#)

Stochastic Gradient Descent, [3](#)

subgradient, [2](#)

Tanh, [4](#)

**List of Theorems**